
Final Project Presentation

-ResNet-

2020095178 최윤선

Colab Link: <https://colab.research.google.com/drive/1osTSRTRdpXRdUTCvDNnAjwDhkjxbx7xv?usp=sharing>

Reference Paper: Deep Residual Learning for Image Recognition (ResNet)

Data Normalization

- For better data normalization, I used the mean and standard deviation of the R, G, B data of the train dataset, instead of 0.5.

```
def normalization(dataset):  
    mean = np.array([np.mean(x.numpy(), axis=(1,2)) for x, _ in dataset])  
    r_mean = mean[:, 0].mean()  
    g_mean = mean[:, 1].mean()  
    b_mean = mean[:, 2].mean()  
  
    std = np.array([np.std(x.numpy(), axis=(1,2)) for x, _ in dataset])  
    r_std = std[:, 0].mean()  
    g_std = std[:, 1].mean()  
    b_std = std[:, 2].mean()  
  
    return (r_mean, g_mean, b_mean), (r_std, g_std, b_std)
```

→ The function of calculating the mean and standard deviation

```
transform = transforms.Compose([transforms.ToTensor()])  
trainset = datasets.CIFAR10(root='./data', train=True, download=True, transform=transform)  
  
mean, std = normalization(trainset)  
print('Mean (R, G, B): ', mean)  
print('Standard deviation (R, G, B): ',std)
```

```
Files already downloaded and verified  
Mean (R, G, B): (0.49139965, 0.48215845, 0.4465309)  
Standard deviation (R, G, B): (0.20220213, 0.20220213, 0.20220213)
```

→ The mean and standard deviation of R, G, B

```
transform = transforms.Compose([transforms.ToTensor(), transforms.Normalize(mean, std)])
```

Data Augmentation

- For higher accuracy, I applied various data augmentation method on train dataset.

```
transform_aug = transforms.Compose([transforms.Resize((32,32)), transforms.RandomHorizontalFlip(),
                                   transforms.RandomRotation(10), transforms.RandomAffine(0, shear=10, scale=(0.8,1.2)),
                                   transforms.ColorJitter(brightness=0.2, contrast=0.2, saturation=0.2),
                                   transforms.ToTensor(), transforms.Normalize(mean, std)])

trainset = datasets.CIFAR10(root='./ data', train=True, download=True, transform=transform_aug)
valset = datasets.CIFAR10(root='./ data', train=True, download=True, transform=transform)
testset = datasets.CIFAR10(root='./ data', train=False, download=True, transform=transform)

classes = ('plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
```

- `transforms.Resize()`: Resize the image.
- `transforms.RandomHorizontalFlip()`: Invert left and right with the defined probability of p. (The default value of p is 0.5.)
- `transforms.RandomRotation()`: Rotate the images randomly at a given angle.
- `transforms.RandomAffine()`: Do a random affine transformation like rotating or moving.
- `transforms.ColorJitter()`: Arbitrarily change brightness, contrast, saturation, color tone.

Data Loader

- Divide the dataset into train, validation, and test dataset.

```
np.random.seed(0)
val_ratio = 0.1
train_size = len(trainset)
indices = list(range(train_size))
split_idx = int(np.floor(val_ratio*train_size))
np.random.shuffle(indices)
train_idx, val_idx = indices[split_idx:], indices[:split_idx]

train_sampler = SubsetRandomSampler(train_idx)
val_sampler = SubsetRandomSampler(val_idx)

batch_size = 128

train_loader = torch.utils.data.DataLoader(trainset, batch_size=batch_size, sampler=train_sampler, num_workers=2)
val_loader = torch.utils.data.DataLoader(valset, batch_size=batch_size, sampler=val_sampler, num_workers=2)
test_loader = torch.utils.data.DataLoader(testset, batch_size=batch_size, shuffle=False, num_workers=2)
```

- Set the batch size to 128 according to the paper.

Model Architecture – ResNet44

- Implement the commonly used convolutional layer as a function.



```
def conv3(in_channels, out_channels, stride=1):  
    return nn.Conv2d(in_channels, out_channels, kernel_size=3, stride=stride, padding=1, bias=False)
```

Model Architecture – ResNet44

- Residual Block

```
class ResidualBlock(nn.Module):
    def __init__(self, in_channels, out_channels, stride=1, shortcut=None):
        super(ResidualBlock, self).__init__()
        self.conv1 = conv3(in_channels, out_channels, stride)
        self.bn1 = nn.BatchNorm2d(out_channels)
        self.relu = nn.ReLU(inplace=True)
        self.conv2 = conv3(out_channels, out_channels)
        self.bn2 = nn.BatchNorm2d(out_channels)
        self.shortcut = shortcut

    def forward(self, x):
        identity = x
        out = self.conv1(x)
        out = self.bn1(out)
        out = self.relu(out)
        out = self.conv2(out)
        out = self.bn2(out)
        if self.shortcut:
            identity = self.shortcut(x)
        out += identity
        out = self.relu(out)
        return out
```



Model Architecture – ResNet44


- Main Block

```
class ResNet(nn.Module):
    def __init__(self, block, layers, num_classes=10):
        super(ResNet, self).__init__()
        self.in_channels = 16
        self.conv = conv3(3, self.in_channels)
        self.bn = nn.BatchNorm2d(self.in_channels)
        self.relu = nn.ReLU(inplace=True)
        self.layer1 = self.make_layer(block, 16, layers[0], 1)
        self.layer2 = self.make_layer(block, 32, layers[1], 2)
        self.layer3 = self.make_layer(block, 64, layers[2], 2)
        self.avgPool = nn.AvgPool2d(8)
        self.fc = nn.Linear(64, num_classes)

    def make_layer(self, block, out_channels, num_blocks, stride=1):
        shortcut = None
        if (stride != 1) or (self.in_channels != out_channels):
            shortcut = nn.Sequential(conv3(self.in_channels, out_channels, stride=stride),
                                     nn.BatchNorm2d(out_channels))

        layers = []
        layers.append(block(self.in_channels, out_channels, stride, shortcut))
        self.in_channels = out_channels
        for _ in range(1, num_blocks):
            layers.append(block(out_channels, out_channels))
        return nn.Sequential(*layers)

    def forward(self, x):
        out = self.conv(x)
        out = self.bn(out)
        out = self.relu(out)
        out = self.layer1(out)
        out = self.layer2(out)
        out = self.layer3(out)
        out = self.avgPool(out)
        out = out.view(out.size(0), -1)
        out = self.fc(out)
        return out
```



output map size	32×32	16×16	8×8
# layers	1+2n	2n	2n
# filters	16	32	64

Model Architecture – ResNet44

- Define the model

```
model = ResNet(ResidualBlock, [7, 7, 7]).to(device)
loss_func = nn.CrossEntropyLoss().to(device)
optimizer = optim.SGD(model.parameters(), lr=0.1, momentum=0.9, weight_decay=1e-4)
scheduler = optim.lr_scheduler.MultiStepLR(optimizer, milestones=[32000, 48000], gamma=0.1)
```

- According to the paper
 - Set $n=7$ and create ResNet model with 44 layers.
 - The optimizer is SGD and set the initial learning rate to 0.1, momentum to 0.9, and weight decay to 0.0001.
 - Use MultiStepLR method to control the learning rate.

Decay the learning rate of each parameter group by $\gamma=0.1$,

once the number of epoch reaches one of the milestones(32000~48000 step).

output map size	32×32	16×16	8×8
# layers	$1+2n$	$2n$	$2n$
# filters	16	32	64

We use a weight decay of 0.0001 and momentum of 0.9, and adopt the weight initialization in [12] and BN [16] but with no dropout. These models are trained with a mini-batch size of 128 on two GPUs. We start with a learning rate of 0.1, divide it by 10 at 32k and 48k iterations, and

Train & Validation

```
epochs = 50
train_loss_list = []
val_loss_list = []
val_acc_list = []

for epoch in range(epochs):
    model.train()
    train_loss = 0
    train_total = 0
    for data, target in train_loader:
        data, target = data.to(device), target.to(device)
        scheduler.step()
        optimizer.zero_grad()
        output = model(data)
        loss = loss_func(output, target)
        loss.backward()
        optimizer.step()
        train_loss += loss.item()
        train_total += target.size(0)

    train_loss /= train_total
    train_loss_list.append(train_loss)

    model.eval()
    correct = 0
    val_total = 0
    val_loss = 0
    with torch.no_grad():
        for data, target in val_loader:
            data, target = data.to(device), target.to(device)
            output = model(data)
            val_loss += loss_func(output, target).item()
            _, prediction = torch.max(output.data, 1)
            val_total += target.size(0)
            correct += prediction.eq(target.view_as(prediction)).sum().item()

    val_loss /= val_total
    val_acc = 100.*correct/val_total

    val_loss_list.append(val_loss)
    val_acc_list.append(val_acc)

    print('Epoch: {} \t Train Loss: {:.4f} \t Valid Loss: {:.4f} \t Valid Accuracy: {:.2f}'.format(epoch+1, train_loss, val_loss, val_acc))



loss_list = torch.tensor(val_loss_list)
acc_list = torch.tensor(val_acc_list)

plt.figure(figsize=(10,5))
plt.subplot(1,2,1)
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.plot(torch.tensor(range(1,51)), train_loss_list, label='train_loss')
plt.plot(torch.tensor(range(1,51)), val_loss_list, label='val_loss')
plt.legend()
plt.subplot(1,2,2)
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.plot(torch.tensor(range(1,51)), val_acc_list)
plt.show()
```

Epoch: 1	Train Loss: 0.015249	Valid Loss: 0.013168	Valid Accuracy: 38.78
Epoch: 2	Train Loss: 0.012389	Valid Loss: 0.011775	Valid Accuracy: 45.82
Epoch: 3	Train Loss: 0.010267	Valid Loss: 0.009481	Valid Accuracy: 58.44
Epoch: 4	Train Loss: 0.008633	Valid Loss: 0.008020	Valid Accuracy: 63.82
Epoch: 5	Train Loss: 0.007490	Valid Loss: 0.007371	Valid Accuracy: 67.60
Epoch: 6	Train Loss: 0.006626	Valid Loss: 0.006011	Valid Accuracy: 73.72
Epoch: 7	Train Loss: 0.006101	Valid Loss: 0.006063	Valid Accuracy: 74.82
Epoch: 8	Train Loss: 0.005592	Valid Loss: 0.005135	Valid Accuracy: 77.76
Epoch: 9	Train Loss: 0.005299	Valid Loss: 0.004809	Valid Accuracy: 79.28
Epoch: 10	Train Loss: 0.004982	Valid Loss: 0.005345	Valid Accuracy: 76.58
Epoch: 11	Train Loss: 0.004821	Valid Loss: 0.004718	Valid Accuracy: 79.64
Epoch: 12	Train Loss: 0.004634	Valid Loss: 0.004646	Valid Accuracy: 80.74
Epoch: 13	Train Loss: 0.004411	Valid Loss: 0.004455	Valid Accuracy: 81.06
Epoch: 14	Train Loss: 0.004271	Valid Loss: 0.004207	Valid Accuracy: 82.36
Epoch: 15	Train Loss: 0.004121	Valid Loss: 0.004892	Valid Accuracy: 79.60
Epoch: 16	Train Loss: 0.003999	Valid Loss: 0.004440	Valid Accuracy: 81.10
Epoch: 17	Train Loss: 0.003854	Valid Loss: 0.003518	Valid Accuracy: 84.66
Epoch: 18	Train Loss: 0.003810	Valid Loss: 0.003982	Valid Accuracy: 82.94
Epoch: 19	Train Loss: 0.003709	Valid Loss: 0.003721	Valid Accuracy: 84.24
Epoch: 20	Train Loss: 0.003596	Valid Loss: 0.003173	Valid Accuracy: 86.10
Epoch: 21	Train Loss: 0.003536	Valid Loss: 0.003499	Valid Accuracy: 85.12
Epoch: 22	Train Loss: 0.003414	Valid Loss: 0.003601	Valid Accuracy: 84.60
Epoch: 23	Train Loss: 0.003375	Valid Loss: 0.003323	Valid Accuracy: 85.58
Epoch: 24	Train Loss: 0.003362	Valid Loss: 0.003459	Valid Accuracy: 85.40
Epoch: 25	Train Loss: 0.003238	Valid Loss: 0.003843	Valid Accuracy: 84.34
Epoch: 26	Train Loss: 0.003221	Valid Loss: 0.004070	Valid Accuracy: 83.14
Epoch: 27	Train Loss: 0.003151	Valid Loss: 0.003843	Valid Accuracy: 84.76
Epoch: 28	Train Loss: 0.003132	Valid Loss: 0.003704	Valid Accuracy: 85.76
Epoch: 29	Train Loss: 0.003063	Valid Loss: 0.003125	Valid Accuracy: 87.02
Epoch: 30	Train Loss: 0.003029	Valid Loss: 0.003445	Valid Accuracy: 85.60
Epoch: 31	Train Loss: 0.003000	Valid Loss: 0.003271	Valid Accuracy: 86.10
Epoch: 32	Train Loss: 0.003013	Valid Loss: 0.003061	Valid Accuracy: 87.10
Epoch: 33	Train Loss: 0.002877	Valid Loss: 0.003114	Valid Accuracy: 87.00
Epoch: 34	Train Loss: 0.002872	Valid Loss: 0.002757	Valid Accuracy: 87.82
Epoch: 35	Train Loss: 0.002846	Valid Loss: 0.002910	Valid Accuracy: 87.18
Epoch: 36	Train Loss: 0.002775	Valid Loss: 0.003052	Valid Accuracy: 87.12
Epoch: 37	Train Loss: 0.002822	Valid Loss: 0.002924	Valid Accuracy: 87.98
Epoch: 38	Train Loss: 0.002759	Valid Loss: 0.003684	Valid Accuracy: 85.66
Epoch: 39	Train Loss: 0.002700	Valid Loss: 0.002963	Valid Accuracy: 87.76
Epoch: 40	Train Loss: 0.002741	Valid Loss: 0.003867	Valid Accuracy: 84.08
Epoch: 41	Train Loss: 0.002665	Valid Loss: 0.003159	Valid Accuracy: 86.96
Epoch: 42	Train Loss: 0.002711	Valid Loss: 0.002969	Valid Accuracy: 87.62
Epoch: 43	Train Loss: 0.002659	Valid Loss: 0.002865	Valid Accuracy: 88.04
Epoch: 44	Train Loss: 0.002546	Valid Loss: 0.003243	Valid Accuracy: 86.76
Epoch: 45	Train Loss: 0.002567	Valid Loss: 0.003175	Valid Accuracy: 87.34
Epoch: 46	Train Loss: 0.002582	Valid Loss: 0.002913	Valid Accuracy: 88.62
Epoch: 47	Train Loss: 0.002531	Valid Loss: 0.002950	Valid Accuracy: 88.08
Epoch: 48	Train Loss: 0.002485	Valid Loss: 0.003898	Valid Accuracy: 84.88
Epoch: 49	Train Loss: 0.002521	Valid Loss: 0.003015	Valid Accuracy: 87.28
Epoch: 50	Train Loss: 0.002511	Valid Loss: 0.002848	Valid Accuracy: 87.52

Train & Validation

- Result Graph



Test – Final Accuracy

```
epochs = 50
for epoch in range(epochs):
    model.eval()
    correct = 0
    total = 0
    with torch.no_grad():
        for data, target in test_loader:
            data, target = data.to(device), target.to(device)
            output = model(data)
            _, prediction = torch.max(output.data, 1)
            total += target.size(0)
            correct += prediction.eq(target.view_as(prediction)).sum().item()

    test_acc = 100.*correct/total
    print('Test Accuracy: {:.2f}'.format(test_acc))
```

Test Accuracy: 87.09

Final accuracy is 87.09%!!